

FlashDeconv: Fast Linear Algebra for Scalable Hybrid Deconvolution

Introduction

This is a fully evolved methodological framework. Based on reflections on Cell2location (accuracy), RCTD (platform effects), and CARD (spatial correlation), this new version addresses the shortcomings of **non-Gaussian distribution handling**, **platform effect correction**, and **rare cell type preservation**, while maintaining the extreme speed brought by **randomized sketching**.

1 Methods: FlashDeconv

1.1 Overview and Statistical Framework

The fundamental challenge in spatial transcriptomics (ST) deconvolution is solving the cell type abundance matrix $\beta \in \mathbb{R}^{N \times K}$ given the spatial gene expression matrix $Y \in \mathbb{R}^{N \times G}$ and the single-cell reference signature matrix $X \in \mathbb{R}^{K \times G}$, such that $Y \approx \beta X$. Current Bayesian methods (e.g., Cell2location) model the raw counts using Negative Binomial distributions, which provides high accuracy but incurs prohibitive computational costs for million-scale datasets. Conversely, fast linear regression methods (e.g., NNLS) often ignore the heteroscedastic noise of count data and platform-specific effects.

FlashDeconv introduces a four-stage framework that combines:

1. **Gene Selection** to identify informative genes (highly variable genes + cell-type markers).
2. **Data Preprocessing** with flexible normalization options (log-CPM, Pearson residuals, or raw).
3. **Structure-Preserving Matrix Sketching** to compress genomic dimensions ($G \approx 20,000$) into a biologically informative low-dimensional subspace ($d \approx 500$), reducing memory usage by orders of magnitude.
4. **Spatial Graph Regularized Optimization** to explicitly model spatial autocorrelation among neighboring spots, solved via a high-performance Coordinate Descent algorithm.

1.2 Data Preprocessing

Raw ST data follows a Poisson-Gamma mixture distribution where variance depends on the mean. Direct application of Euclidean-based sketching (Johnson-Lindenstrauss lemma) on raw counts is statistically invalid. FlashDeconv provides three preprocessing strategies, selectable via the `preprocess` parameter:

1.2.1 Method 1: Log-CPM Normalization (Default)

The default and recommended preprocessing (`preprocess="log_cpm"`) normalizes counts to Counts Per Million (CPM) and applies log1p transformation:

$$\tilde{Y}_{ig} = \log(1 + 10^4 \cdot \frac{Y_{ig}}{N_i}), \quad \tilde{X}_{kg} = \log(1 + 10^4 \cdot \frac{X_{kg}}{N_k^{ref}}) \quad (1)$$

where $N_i = \sum_g Y_{ig}$ is the total count for spot i and $N_k^{ref} = \sum_g X_{kg}$ is the total count for cell type k . This normalization:

- Removes sequencing depth effects
- Stabilizes variance via log transformation
- Preserves non-negativity for downstream NNLS

1.2.2 Method 2: Uncentered Pearson Residuals

For variance-stabilizing transformation (`preprocess="pearson"`), we use **uncentered** Pearson residuals that divide by the expected standard deviation:

$$\tilde{Y}_{ig} = \frac{Y_{ig}}{\sigma_g^Y}, \quad \tilde{X}_{kg} = \frac{X_{kg}}{\sigma_g^X} \quad (2)$$

where the gene-wise standard deviations are computed using a Negative Binomial variance model:

$$\sigma_g^Y = \sqrt{\bar{\mu}_g^Y + (\bar{\mu}_g^Y)^2/\theta}, \quad \sigma_g^X = \sqrt{\bar{\mu}_g^X + (\bar{\mu}_g^X)^2/\theta} \quad (3)$$

Here $\bar{\mu}_g^Y = \text{mean}_i(Y_{ig})$ and $\bar{\mu}_g^X = \text{mean}_k(X_{kg})$ are the gene-wise means, and $\theta = 100$ is a fixed overdispersion parameter.

Key Design Choice: Unlike standard Pearson residuals $(Y - \mu)/\sigma$, we use uncentered residuals Y/σ to preserve non-negativity. This is essential because downstream solving uses non-negative least squares (NNLS), which requires non-negative inputs.

1.2.3 Method 3: Raw Counts

For pre-normalized data or synthetic experiments (`preprocess="raw"`), the raw counts can be used directly:

$$\tilde{Y} = Y, \quad \tilde{X} = X \quad (4)$$

Recommendation: Use `log_cpm` (default) for most real datasets. Use `pearson` for data with high technical noise. Use `raw` only for pre-normalized or synthetic data.

1.3 Structure-Preserving Randomized Sketching

Solving the regression problem in the original G -dimensional space ($G \approx 20,000$) is computationally wasteful due to the high collinearity of gene expression. We propose **Structure-Preserving Sketching**, an improvement over standard random projections.

We define a sketching matrix $\Omega \in \mathbb{R}^{G \times d}$, where $d \ll G$ (default $d = 512$). Instead of a dense Gaussian matrix (which is slow to generate and multiply), we use a **Sparse CountSketch Matrix** with an **Importance Sampling** twist:

1. **Feature Selection:** We select a subset of Informative Genes \mathcal{G}_{info} (Union of Highly Variable Genes and Cell-type Specific Markers) to prevent noise accumulation.
2. **Weighted Sketching:** For each column $j \in \{1 \dots d\}$ of Ω , we assign exactly one non-zero value ± 1 for each gene g , but the probability of assignment is weighted by the gene's **Leverage Score** (biological importance). This ensures that marker genes for rare cell types are preserved with high probability.

1.3.1 CountSketch Matrix Construction (Detailed Algorithm)

The CountSketch matrix $\Omega \in \mathbb{R}^{G \times d}$ is constructed as follows:

Input: Number of genes G , sketch dimension d , leverage scores $\ell \in \mathbb{R}^G$

Algorithm:

1. Normalize leverage scores: $p_g = \ell_g / \sum_{g'} \ell_{g'}$ (probability distribution)
2. For each gene $g \in \{1, \dots, G\}$:

- (a) Hash assignment: $h(g) \sim \text{Uniform}(\{1, \dots, d\})$
 - (b) Sign assignment: $s(g) \sim \text{Uniform}(\{-1, +1\})$
 - (c) Importance weight: $w_g = \sqrt{p_g \cdot G}$ (scaled by leverage)
 - (d) Clip weight: $w_g = \text{clip}(w_g, 0.1, 10.0)$ (numerical stability)
3. Construct sparse matrix: $\Omega[g, h(g)] = s(g) \cdot w_g$
4. Column normalization: For each column j :

$$\Omega[:, j] \leftarrow \Omega[:, j] \cdot \frac{\sqrt{G/d}}{\|\Omega[:, j]\|_2} \quad (5)$$

Output: Sparse matrix Ω of shape (G, d) with exactly G non-zero entries

The key innovation is the **importance weighting** $w_g = \sqrt{p_g \cdot G}$: genes with higher leverage scores contribute more strongly to the sketch, preserving discriminative information for rare cell types.

1.3.2 Leverage Score Computation

Leverage scores quantify the importance of each gene for distinguishing cell types. We compute them via SVD of the reference matrix:

$$X^T = U\Sigma V^T \quad (6)$$

where $X \in \mathbb{R}^{K \times G}$ is the (centered) reference signature matrix.

The leverage score for gene g is:

$$\ell_g = \sum_{j=1}^r \frac{\sigma_j^2}{\sigma_j^2 + \epsilon} \cdot U_{gj}^2 \quad (7)$$

where $r = \min(K, G)$ is the rank, σ_j are singular values, and $\epsilon = 10^{-6}$ prevents division by zero. This weighted sum emphasizes genes that lie in the principal subspace of the cell type signatures.

Intuition: Genes with high leverage scores are those that:

- Have high variance across cell types (captured by large σ_j)
- Are not redundant with other genes (unique directions in U)
- Are essential for distinguishing rare cell types (marker genes)

1.3.3 Integration: CountSketch + Importance Sampling

The combination works as follows:

Component	Standard CountSketch	Our Method
Hash function	Uniform random	Uniform random
Sign function	± 1 uniform	± 1 uniform
Entry value	± 1	$\pm w_g$ (leverage-weighted)
Expectation	$\mathbb{E}[\Omega^T \Omega] = I$	$\mathbb{E}[\Omega^T \Omega] \approx I$ (weighted)

Theoretical Guarantee (Leverage Score Sampling Framework):

Following the leverage score sampling framework (Drineas et al., 2006; Mahoney, 2011), the weighted sketch preserves the column space of X with high probability. Specifically, for the regression problem $\min_{\beta} \|Y - \beta X\|_F^2$, solving in the sketched space yields a $(1 + \epsilon)$ -approximation:

$$\|\tilde{Y} - \hat{\beta} \tilde{X}\|_F^2 \leq (1 + \epsilon) \min_{\beta} \|\tilde{Y} - \beta \tilde{X}\|_F^2 \quad (8)$$

with probability at least $1 - \delta$, where $d = O(K \log(K)/\epsilon^2)$ suffices.

The importance weighting ensures that the effective approximation error is **lower for high-leverage genes** (cell type markers), preserving rare cell type signals even after dimension reduction. Unlike standard Johnson-Lindenstrauss projections, this leverage-based approach provides **non-uniform preservation** tailored to the regression structure.

We then project the data into the low-dimensional “sketch space”:

$$Y_{\text{sketch}} = \tilde{Y}\Omega \quad (N \times d) \quad (9)$$

$$X_{\text{sketch}} = \tilde{X}\Omega \quad (K \times d) \quad (10)$$

Theoretical Guarantee: According to the Johnson-Lindenstrauss lemma and recent extensions for oblivious subspace embeddings, solving the regression in this d -dimensional sketch space yields a solution $\hat{\beta}$ that is an $(1 + \epsilon)$ -approximation of the optimal solution in the original space, with high probability.

1.4 Spatial Graph Laplacian Regularization

To incorporate spatial information (inspired by CARD) without the computational burden of inverting dense covariance matrices, we introduce a **Graph Laplacian Regularizer**.

We construct a spatial neighbor graph adjacency matrix A , where $A_{ij} = 1$ if spot i and j are physical neighbors, else 0. The Laplacian matrix is $L = D - A$, where D is the degree matrix.

Our final objective function in the sketch space is:

$$\mathcal{L}(\beta) = \underbrace{\frac{1}{2} \|Y_{\text{sketch}} - \beta X_{\text{sketch}}\|_F^2}_{\text{Fidelity in Sketch Domain}} + \underbrace{\lambda \cdot \text{Tr}(\beta^T L \beta)}_{\text{Spatial Smoothing}} + \underbrace{\rho \|\beta\|_1}_{\text{Sparsity}} \quad (11)$$

Subject to: $\beta \geq 0$.

The term $\text{Tr}(\beta^T L \beta) = \sum_{(i,j) \in E} \|\beta_i - \beta_j\|^2$ forces neighboring spots to have similar cell type compositions, effectively denoising “drop-out” events in low-coverage spots.

1.5 Fast Optimization: Block Coordinate Descent

While the objective function is convex, standard gradient descent is slow for large N . We derive a **Block Coordinate Descent (BCD)** algorithm that solves row-by-row (spot-by-spot) but leverages shared pre-computed matrices.

For each spot i , the update rule for cell type k (β_{ik}) has a **closed-form solution** combining soft-thresholding (for L1) and projection to non-negative (for constraints).

Let $G = X_{\text{sketch}} X_{\text{sketch}}^T$ (the Gram matrix, precomputed once). The update is:

$$r_{ik} = (Y_{\text{sketch}})_i (X_{\text{sketch}})_k^T - \sum_{j \neq k} \beta_{ij} G_{jk} + \lambda \sum_{n \in \mathcal{N}(i)} \beta_{nk} \quad (12)$$

$$\beta_{ik} \leftarrow \left[\frac{S_\rho(r_{ik})}{G_{kk} + \lambda |\mathcal{N}(i)|} \right]_+ \quad (13)$$

where $S_\rho(x) = \text{sign}(x) \cdot \max(0, |x| - \rho)$ is the **soft-thresholding operator** for L1 regularization, and $[\cdot]_+ = \max(0, \cdot)$ enforces non-negativity.

Convergence Criteria:

- **Maximum iterations:** 100 (default)
- **Tolerance:** $\text{rel_change} = \frac{\max |\beta^{(t)} - \beta^{(t-1)}|}{\max |\beta^{(t-1)}| + 10^{-10}} < 10^{-4}$
- Empirically, convergence is achieved in **10-30 iterations** for most datasets.

Computational Edge:

- The term $(X_{\text{sketch}})^T X_{\text{sketch}}$ is a tiny $K \times K$ matrix computed only once.
- The neighbor sum $\sum_{n \in \mathcal{N}(i)} \beta_{nk}$ exploits the sparsity of the spatial graph.
- The algorithm converges in very few iterations and is fully parallelizable across spots.

1.6 Complexity Analysis

Component	Cell2location (VI)	CARD (CAR)	FlashDeconv (Ours)
Data Representation	Dense $N \times G$	Dense $N \times G$	Sketch $N \times d$
Spatial Model	N/A (or slow MRF)	Dense Matrix Inversion	Sparse Laplacian
Time Complexity	$O(\text{Iter} \cdot N \cdot G \cdot K)$	$O(N^3)$ (approx)	$O(\text{Iter} \cdot N \cdot d \cdot K)$
Memory Usage	High (GBs to TBs)	High ($N \times N$ matrix)	Ultra-Low ($O(Nd)$)
1M Spots Time	> 24 Hours (GPU)	Out of Memory	~ 10 Minutes (CPU)

Table 1: Computational complexity comparison of different deconvolution methods.

1.7 Implementation Details for Reproducibility

- **Software:** Implemented in Python using `Numba` for JIT compilation of the BCD solver.
- **Hardware Independence:** Unlike Tangram or Cell2location, FlashDeconv requires **no GPU** and runs efficiently on standard laptops (e.g., MacBook Air with 8GB RAM).

Default Hyperparameters:

Parameter	Default	Valid Range	Description
d (sketch_dim)	512	[128, 2048]	Sketch dimension
λ (lambda_spatial)	5000 or "auto"	$[0, \infty)$	Spatial smoothing; 0 disables
ρ (rho_sparsity)	0.01	[0, 0.1]	L1 regularization for sparsity
n_hvg	2000	[500, 5000]	Number of highly variable genes
n_markers_per_type	50	[20, 200]	Marker genes per cell type
k_neighbors	6	[4, 20]	Neighbors for spatial graph
preprocess	"log_cpm"	{log_cpm, pearson, raw}	Preprocessing method
max_iter	100	—	Maximum BCD iterations
tol	10^{-4}	—	Convergence tolerance

Automatic λ Tuning:

When `lambda_spatial="auto"`, we estimate λ by balancing the data fidelity term and the spatial regularization term. The key insight is that λ must be scaled relative to the Gram matrix $X_{\text{sketch}}^T X_{\text{sketch}}$ to have meaningful effect on the optimization.

In the BCD update, the denominator is:

$$\text{denom} = (X_{\text{sketch}}^T X_{\text{sketch}})_{kk} + \lambda \cdot n_{\text{neighbors}} \quad (14)$$

For spatial regularization to contribute meaningfully, we set:

$$\lambda = \alpha \cdot \frac{\text{mean}(\text{diag}(X_{\text{sketch}}^T X_{\text{sketch}}))}{\bar{d}} \quad (15)$$

where:

- $\alpha = 0.005$ (default) controls the relative strength of spatial regularization
- \bar{d} is the average number of neighbors per spot

This ensures that the spatial term contributes approximately $\alpha/(1 + \alpha) \approx 0.5\%$ to the Hessian diagonal. Users can adjust the effective regularization strength by setting a fixed `lambda_spatial` value (recommended range: 1000-10000 for Visium data).

Numerical Stability Clipping:

Quantity	Clip Range	Rationale
Sketch weights	[0.1, 10]	Balanced importance sampling
Leverage scores	Normalized to sum=1	Proper probability distribution

Key Design Decisions

1. **Flexible Preprocessing:** Unlike Cell2location/RCTD which require specific data formats, FlashDeconv supports multiple preprocessing modes (`log_cpm`, `pearson`, `raw`) to accommodate different data types and user preferences.
2. **Spatial Regularization via Graph Laplacian:** Inspired by CARD’s spatial modeling, but using sparse Graph Laplacian instead of dense CAR matrix inversion. This provides $O(N \cdot K)$ complexity instead of $O(N^3)$.
3. **Structure-Preserving Sketching:** Leverage-weighted CountSketch ensures rare cell type markers are preserved during dimension reduction, addressing the “random projection loses rare cells” concern.
4. **CPU-only, Memory-efficient:** No GPU required. The algorithm runs efficiently on standard laptops (8GB RAM), making it accessible for Stereo-seq/Xenium users with million-scale datasets.
5. **Parameter Sensitivity:** The `lambda_spatial` parameter may require tuning for different data types:
 - For sequencing-based platforms (10x Visium): default 5000 works well
 - For imaging-based platforms (seqFISH+, MERFISH): consider `lambda_spatial=0` to disable spatial regularization when sample sizes are small